

EN 601.473/601.673: Cognitive Artificial Intelligence (CogAI)



**Lecture 8:
Intro to Gen,
importance sampling**

Tianmin Shu

Common questions about Pset 1

- Relevant lecture notes: Lecture 4
- Hypothesis averaging for prediction


Hypothesis averaging:

Compute the probability that C applies to some new object y (i.e., y is a yes number) by averaging the predictions of all hypotheses h , weighted by $p(h|X)$:

$$p(y \in C \mid X) = \sum_{h \in H} \underbrace{p(y \in C \mid h)}_{= \begin{cases} 1 & \text{if } y \in h \\ 0 & \text{if } y \notin h \end{cases}} p(h \mid X)$$

Step 1: Concept inference $p(h \mid X)$

Examples:
16





Step 2: Prediction

For a new number y :
64

$$p(y \in C | X) = \sum_{h \in H} \underbrace{p(y \in C | h)}_{\substack{= 1 \text{ if } y \in h \\ = 0 \text{ if } y \notin h}} p(h | X)$$

Weighted average of 0s and 1s

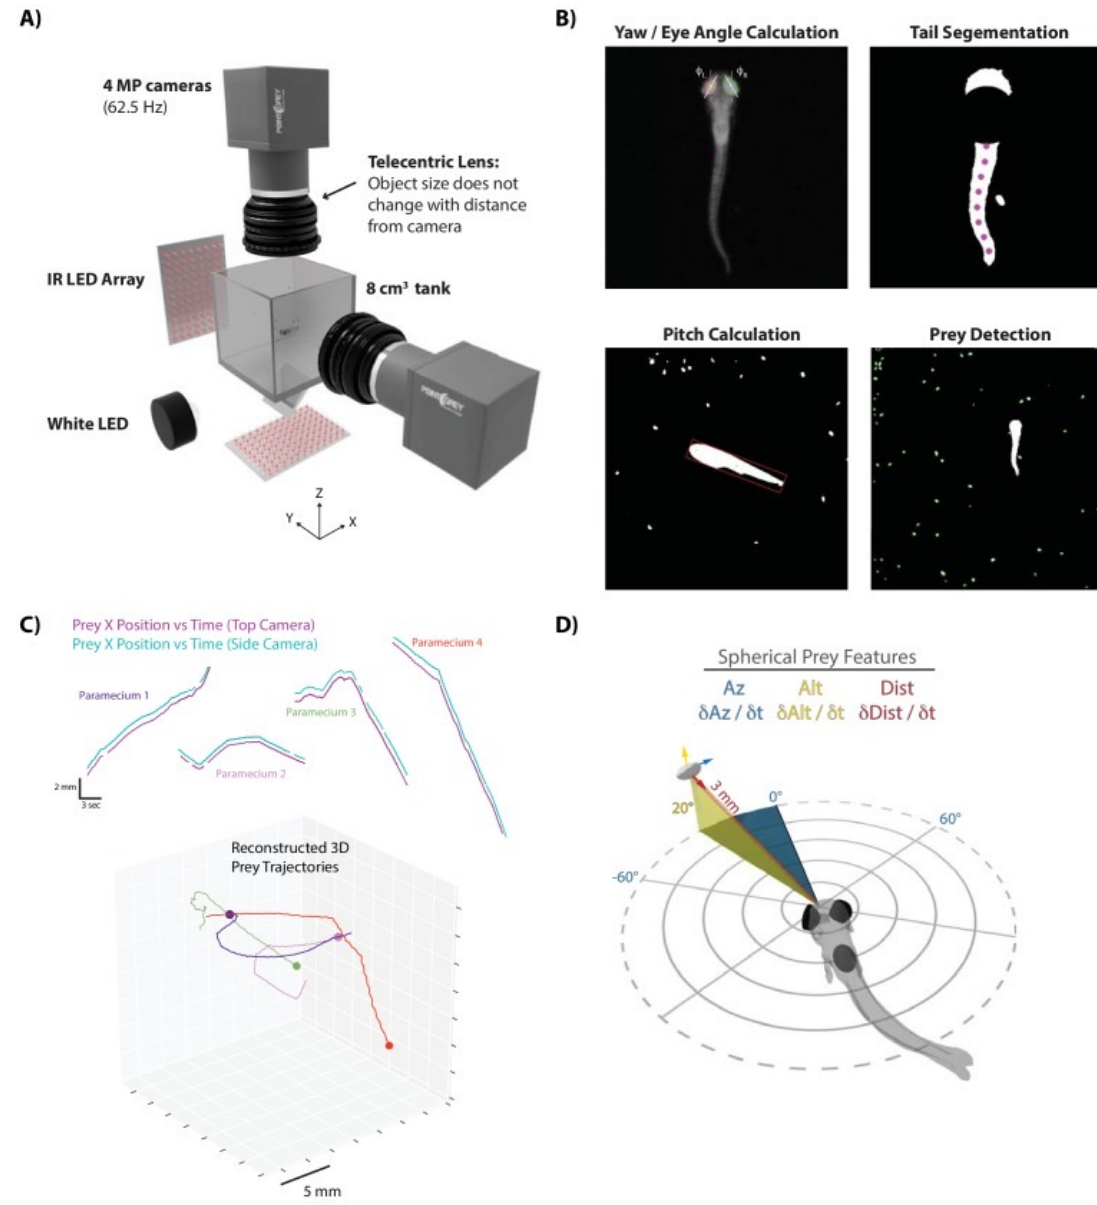


Introduction to Gen

- Gen vs Conventional PPLs
- “Hello world” in Gen: write a simple generative program
- Trace, weights
- Importance sampling & importance resampling

Probabilistic programming in Science

Experimental Neuroscience




Bolton et al. (2019) using
BayesDB

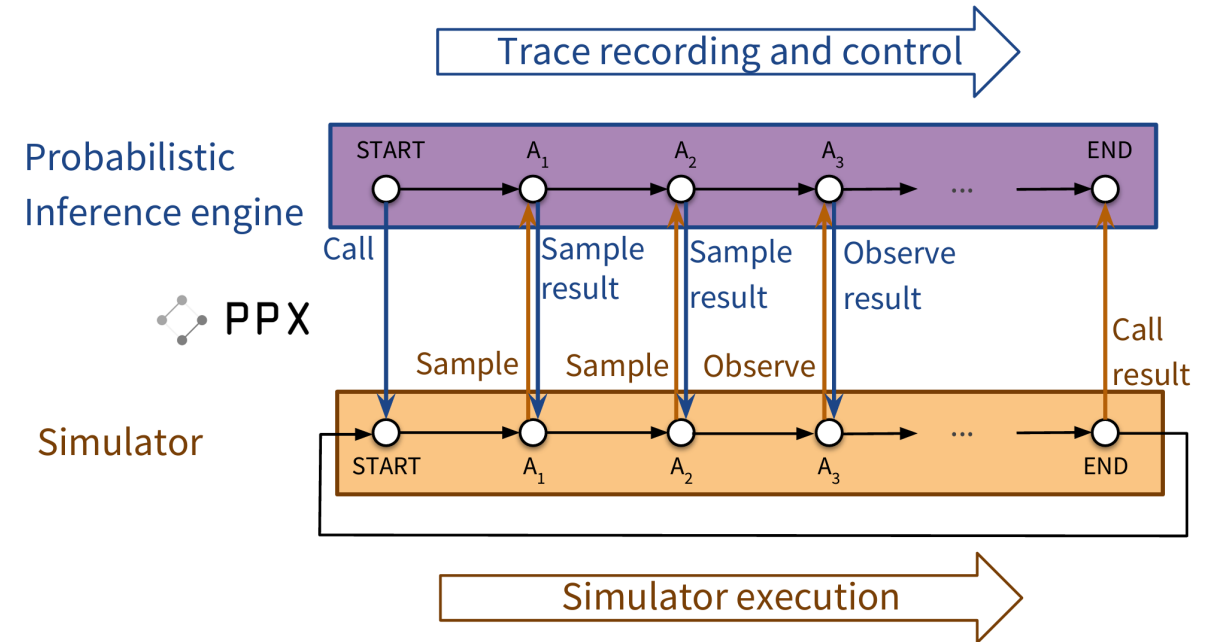
Probabilistic programming in Science

Particle Physics

Determine the properties of particles at the Large Hadron Collider (LHC) at CERN



Bolton et al. (2019) using **pyprob (PyTorch-based PPL) + Sherpa (C++ Simulator)**



Probabilistic programming in Science

Epidemiology

Estimating the number of infections and the impact of
non-pharmaceutical interventions on COVID-19 in European
countries: technical description update

Seth Flaxman*, Swapnil Mishra*, Axel Gandy*, H Juliette T Unwin, Helen Coupland,
Thomas A Mellan, Harrison Zhu, Tresnia Berah, Jeffrey W Eaton, Pablo N P Guzman, Nora
Schmit, Lucia Callizo, Imperial College COVID-19 Response Team, Charles Whittaker, Peter
Winskill, Xiaoyue Xi, Azra Ghani, Christl A. Donnelly, Steven Riley, Lucy C Okell, Michaela
A C Vollmer, Neil M. Ferguson and Samir Bhatt*,¹

Using Stan PPL

Two types of PPLs

- Write your generative programs and let a black box engine run the inference
- Programmable generative models and programmable inference



Automatic Inference

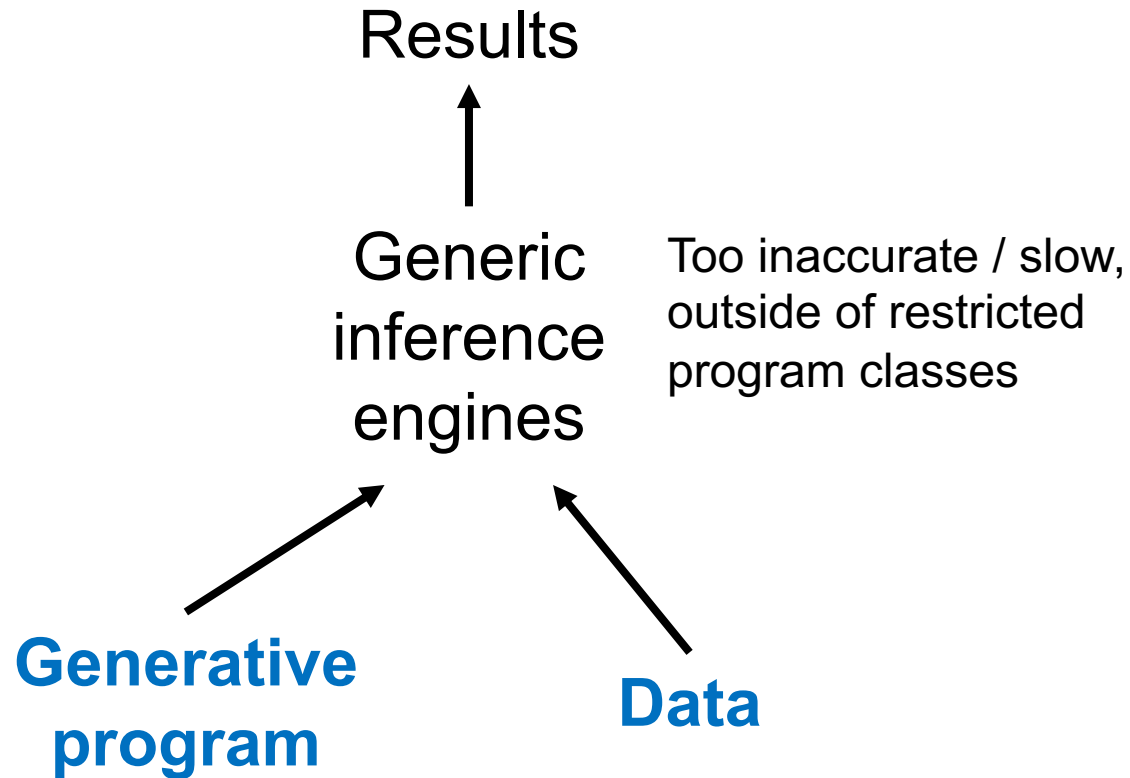


Programmable Inference

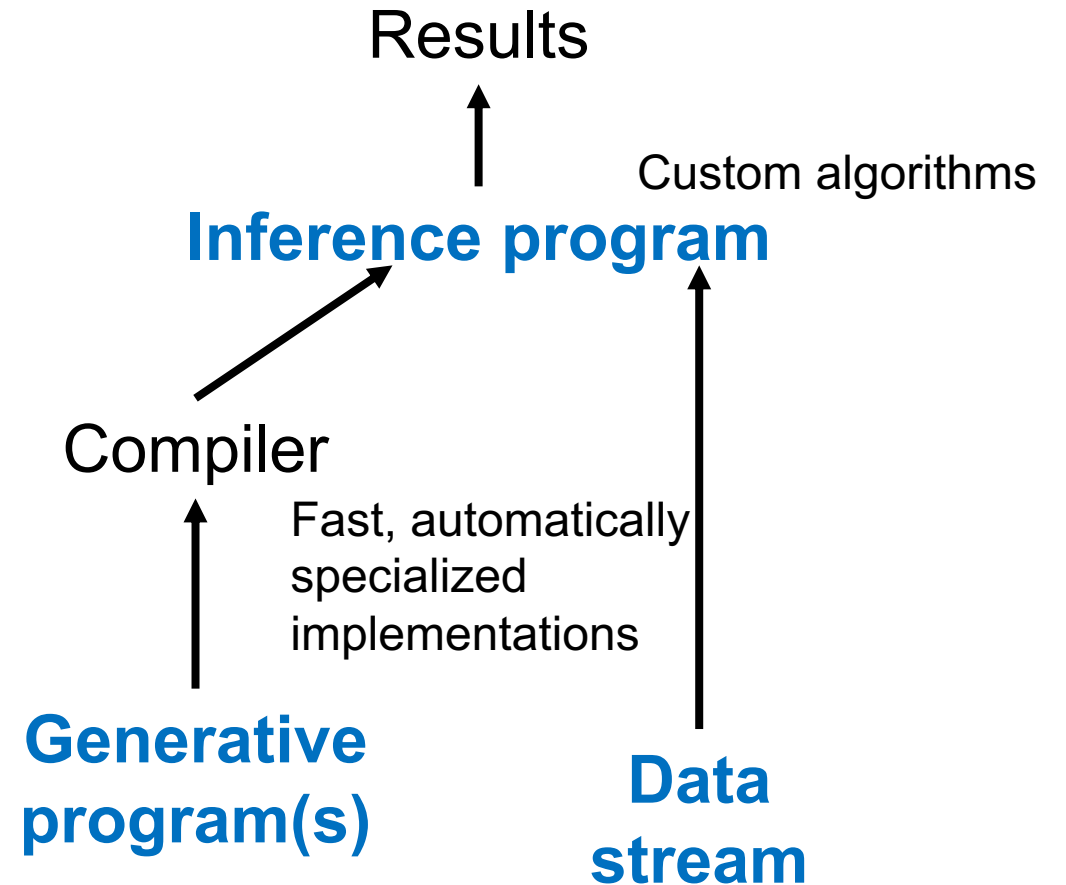
Two types of PPLs

(BLUE: specified by users)

Traditional PPLs



Gen



Competitive performance compared against restricted PPLs

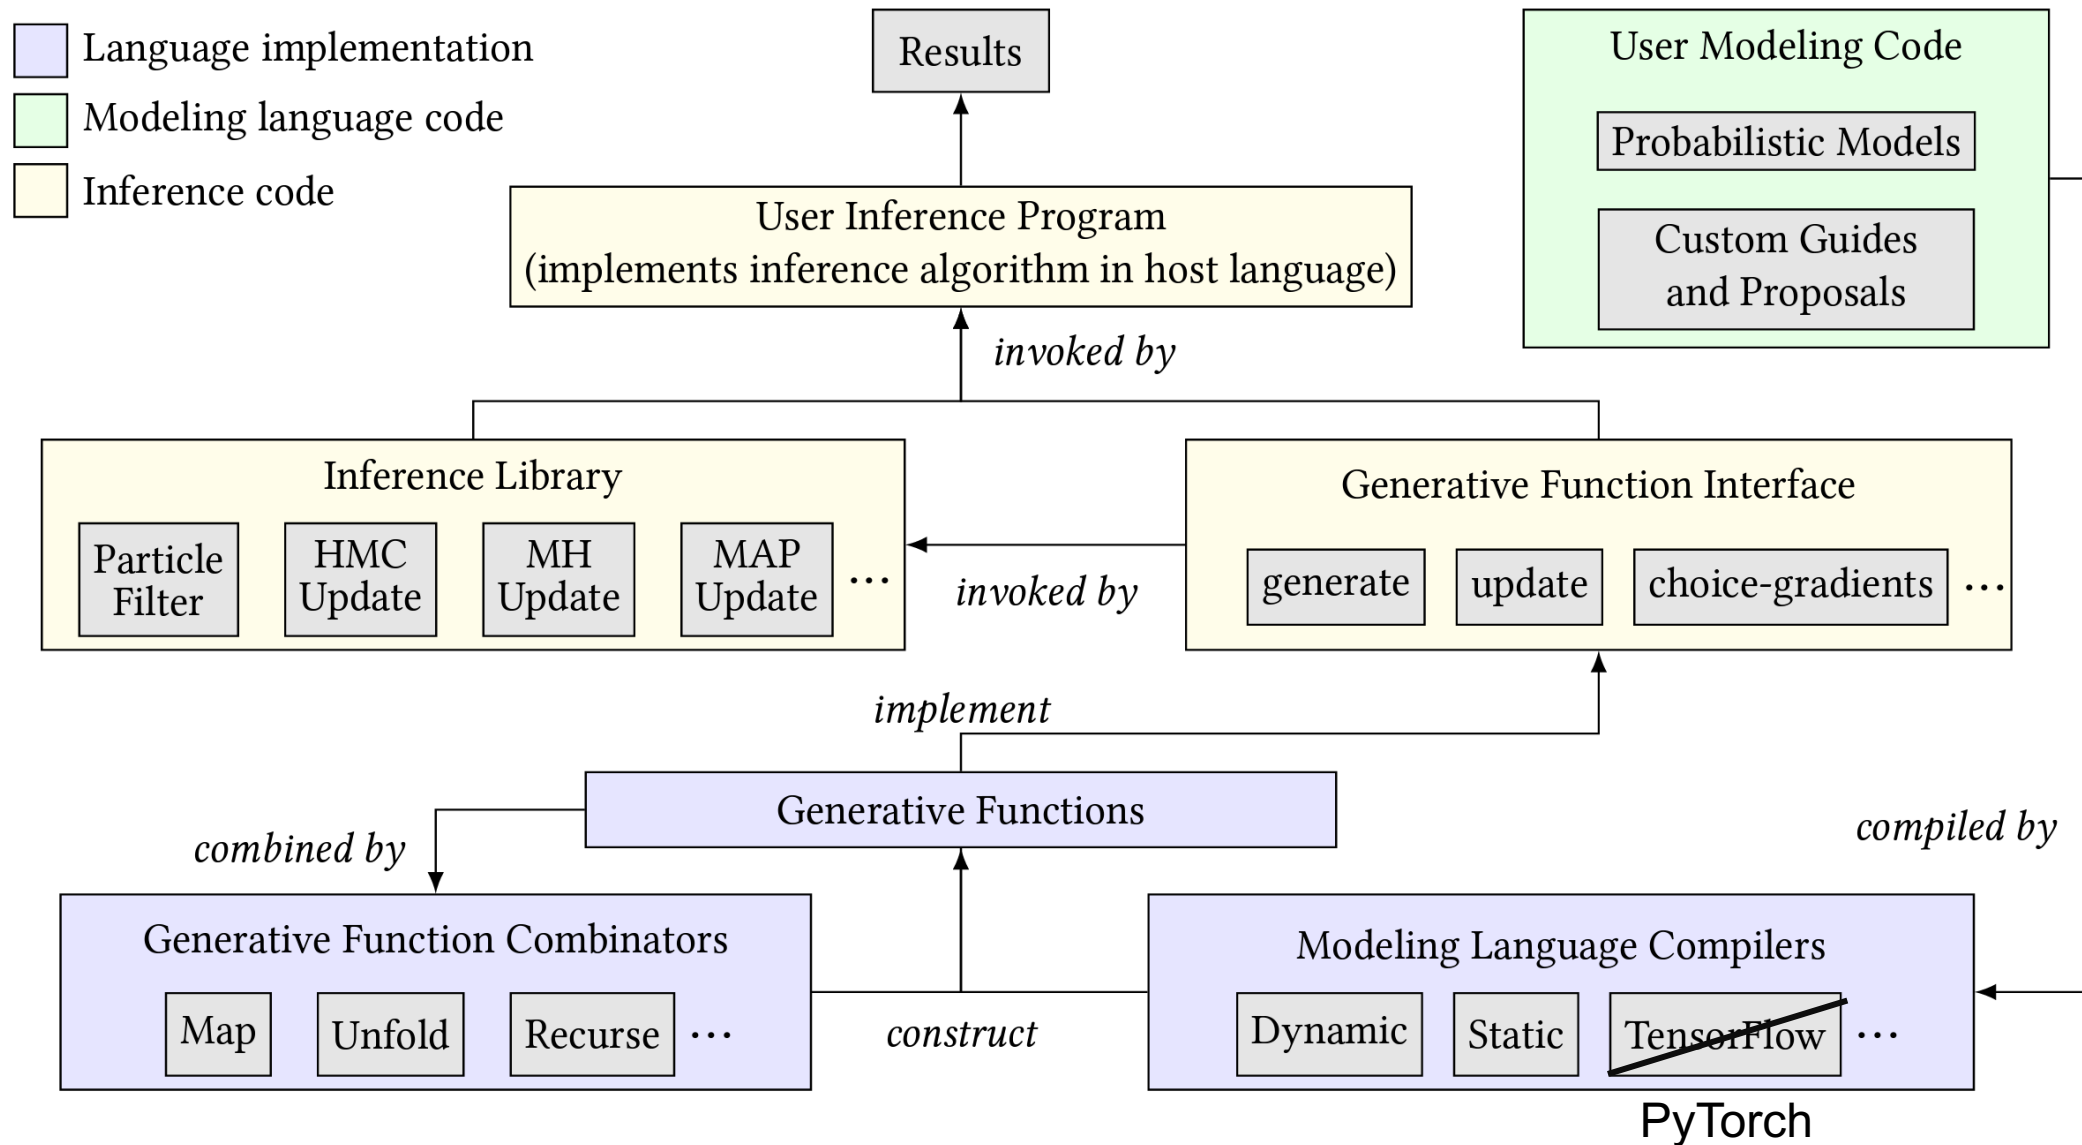
	Inference Algorithm	Runtime (ms)
Stan	Hamiltonian Monte Carlo (NUTS)	53.4ms
Gen (SML + Map)	Gaussian Drift Metropolis Hastings	75.3ms
Edward	Hamiltonian Monte Carlo	76.6ms
Anglican	Gaussian Drift Metropolis Hastings	783ms
Venture	Gaussian Drift Metropolis Hastings	1.3×10^6 ms

Bayesian linear regression

	Proposal Distribution	Runtime (ms)
Gen (DML + Unfold)	Custom	4.9ms (± 0.07)
Gen (DML + Unfold)	Generic	82ms (± 3.6)
Anglican	Generic	275ms (± 11)
Turing	Generic	1174ms (± 25)
Venture	Generic	$> 10^6$ ms

Nonlinear state-space model

Gen's architecture



An example generative model in Gen

Defining a generative model in Julia

```
using Gen: uniform_discrete, bernoulli, categorical

function f(p)
    n = uniform_discrete(1, 10)
    if bernoulli(p)
        n *= 2
    end
    return categorical([i == n ? 0.5 : 0.5/19 for i=1:20])
end;
```

- Sample n uniformly from 1-10
- With prob of p, multiply n by 2
- With 0.5, sample n, and with 0.5 sample uniformly from the remaining 19 numbers in 1-20

@gen macro for defining a generative model in Gen

```
using Gen: @gen

@gen function gen_f(p)
    n = { :initial_n } ~ uniform_discrete(1, 10)
    if ( { :do_branch } ~ bernoulli(p) )
        n *= 2
    end
    return { :result } ~ categorical([i == n ? 0.5 : 0.5/19 for i=1:20])
end;
```

```
var = { :address } ~ distribution(args...)
address ~ distribution(args...)
```

Trace

```
using Gen: @gen

@gen function gen_f(p)
    n = {:initial_n} ~ uniform_discrete(1, 10)
    if ( {:do_branch} ~ bernoulli(p))
        n *= 2
    end
    return {:result} ~ categorical([i == n ? 0.5 : 0.5/19 for i=1:20])
end;
```


```
trace = simulate(gen_f, (0.3,));
```

```
get_choices(trace)
```

```
├── :result : 19
├── :do_branch : true
└── :initial_n : 7
```

Generating sequence

- Example: Unknown urn problem
- Determine the ratio of the balls based on observations




```
@gen function unknown_urn()  
  #  $p(\theta) \sim \text{Uniform}(0,1)$  [prior distribution]  
  theta ~ uniform(0, 1)  
  for i=1:100  
    #  $p(y=1|\theta) \sim \text{Bernoulli}(\theta)$  [likelihood function]  
    { :data => i => :y } ~ bernoulli(theta)  
  end  
end
```

```
(trace, _) = generate(unknown_urn_static, (3,))  
get_choices(trace)
```

```
|  
├── :theta : 0.48592872586107905  
└── :data  
    ├── 1  
    │   └── :y : false  
    ├── 2  
    │   └── :y : true  
    └── 3  
        └── :y : false
```

Importance sampling

- **Goal:** approximate a *target* distribution $P(x)$, which is hard to sample from
- **Main idea:** Sample a set of particles from a simple *proposal* distribution $q(x)$ (e.g., a uniform distribution) and *weight* them to approximate the distribution
- For each particle i
 - Sample: $x_i \sim q(x)$
 - Weight: $w_i = \frac{p(x_i)}{q(x_i)}$
 - (x_i, w_i)




Importance resampling

- **Goal:** sample from the approximated distribution based on the particles and their weights
- Normalize the weights of particles (probability of each particle assuming the whole population is the sampled particles):

$$w_i = \frac{w_i}{\sum_j w_j}$$

- For each new sample x ,
 - Sample $j \in \{1, 2, \dots, K\}$, from with probability w_j
 - The new sample $x = x_j$
 - Set weight $w_i = 1/K$

Limit: proposal distribution needs to be reasonably close to the target distribution



Importance resampling in Bayesian inference

- Target: $h_i \sim P(h|D)$
- Sample: $h_i \sim Q(h)$
- Weight: $w_i \propto \frac{P(D|h_i)P(h_i)}{Q(h_i)}$ $w_i = \frac{w_i}{\sum_j w_j}$
- (h_i, w_i)
- For each new sample h ,
 - Sample $j \in \{1, 2, \dots, K\}$, from with probability w_j
 - The new sample $h = h_j$

Sample, retrieve, reject, weight, update a trace

See the jupyter notebook