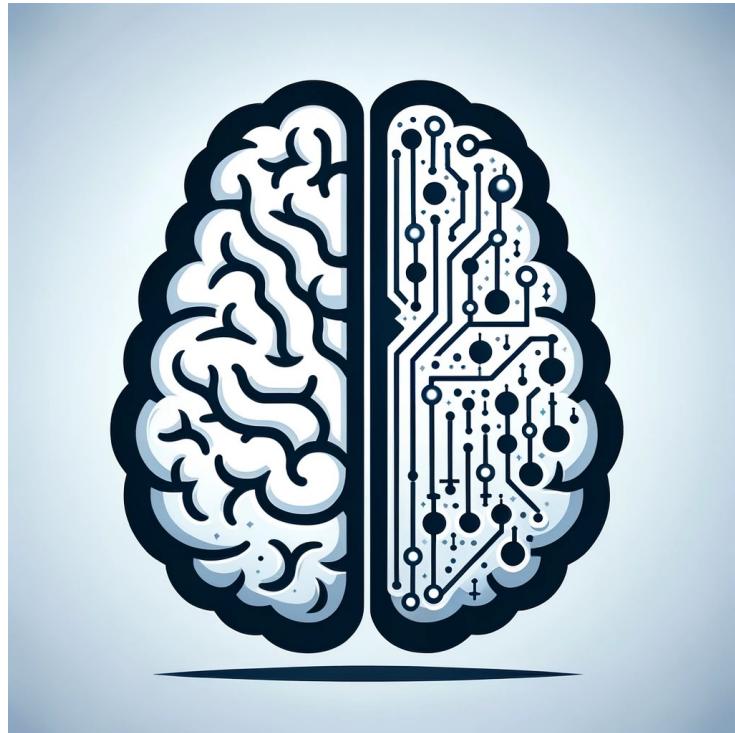


EN 601.473/601.673: Cognitive Artificial Intelligence (CogAI)



**Lecture 11:
MCMC in Gen,
RJMCMC, SMC**

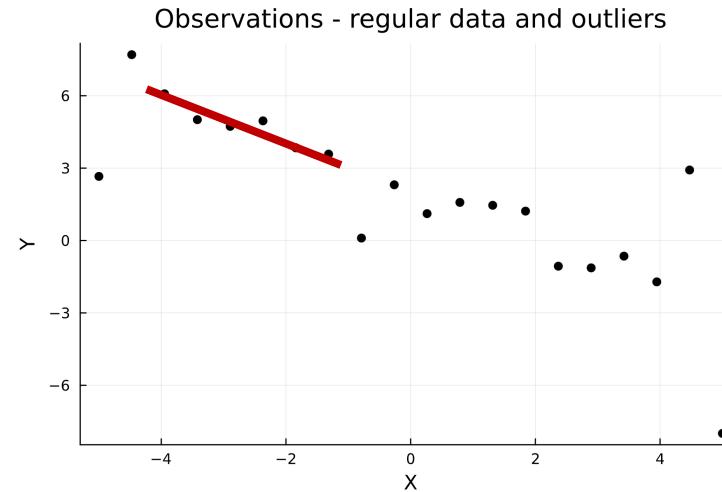
Tianmin Shu

Problem set 2 release

- Release: next lecture, Feb 29
- Due by March 26, end of day
- Q1: importance sampling vs MCMC in Gen
- Q2: SMC (particle filtering) in Gen
- Start as early as possible (also, go through the notebooks for the lectures)

Use heuristics to make smarter proposals

- Data-driven proposals $q(H | \text{data})$



- An example: RANSAC (Random Sample Consensus)
- Repeatedly choose a small random subset of the points, say, of size 3.
- Conduct least-squares linear regression to find a line of best fit for those points.
- Count how many points (from the entire set) are near the line we found.
- After a suitable number of iterations (say, 10), return the line that had the highest score.

RANSAC in Julia

```
import StatsBase

struct RANSACParams
    """the number of random subsets to try"""
    iters::Int

    """the number of points to use to construct a hypothesis"""
    subset_size::Int

    """the error threshold below which a datum is considered an inlier"""
    eps::Float64

    function RANSACParams(iters, subset_size, eps)
        if iters < 1
            error("iters < 1")
        end
        new(iters, subset_size, eps)
    end
end
```

```

function ransac(xs::Vector{Float64}, ys::Vector{Float64}, params::RANSACParams)
    best_num_inliers::Int = -1
    best_slope::Float64 = NaN
    best_intercept::Float64 = NaN
    for i=1:params.iters
        # select a random subset of points
        rand_ind = StatsBase.sample(1:length(xs), params.subset_size, replace=false)
        subset_xs = xs[rand_ind]
        subset_ys = ys[rand_ind]

        # estimate slope and intercept using least squares
        A = hcat(subset_xs, ones(length(subset_xs)))
        slope, intercept = A \ subset_ys # use backslash operator for least sq soln

        ypred = intercept .+ slope * xs

        # count the number of inliers for this (slope, intercept) hypothesis
        inliers = abs.(ys - ypred) .< params.eps
        num_inliers = sum(inliers)

        if num_inliers > best_num_inliers
            best_slope, best_intercept = slope, intercept
            best_num_inliers = num_inliers
        end
    end

    # return the hypothesis that resulted in the most inliers
    (best_slope, best_intercept)
end;

```

RANSAC proposal in Gen

- Wrap the ransac Julia function in a Gen proposal

```
@gen function ransac_proposal(prev_trace, xs, ys)
    (slope_guess, intercept_guess) = ransac(xs, ys, RANSACParams(10, 3, 1.))
    slope ~ normal(slope_guess, 0.1)
    intercept ~ normal(intercept_guess, 1.0)
end;
```

One iteration of RANSAC-based update

```
function ransac_update(tr)
    # Use RANSAC to (potentially) jump to a better line
    # from wherever we are
    (tr, _) = mh(tr, ransac_proposal, (xs, ys))

    # Spend a while refining the parameters, using Gaussian drift
    # to tune the slope and intercept, and resimulation for the noise
    # and outliers.
    for j=1:20
        (tr, _) = mh(tr, select(:prob_outlier))
        (tr, _) = mh(tr, select(:noise))
        (tr, _) = mh(tr, line_proposal, ())
        # Reclassify outliers
        for i=1:length(get_args(tr)[1])
            (tr, _) = mh(tr, select(:data => i => :is_outlier))
        end
    end
    tr
end
```

RANSAC-based MCMC inference

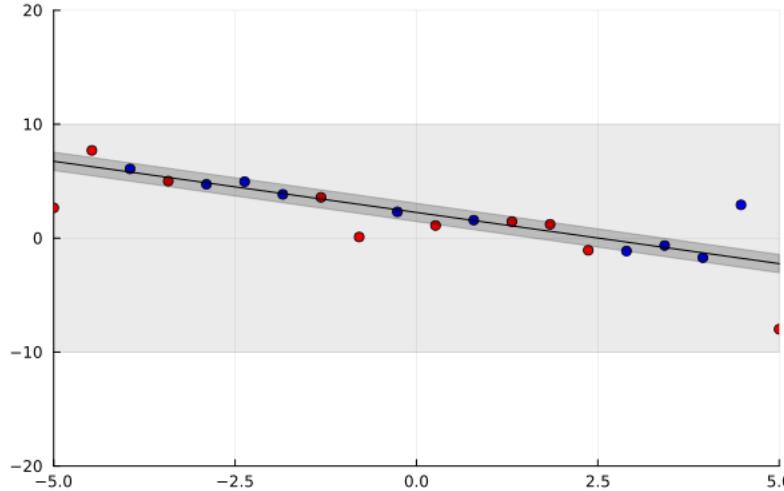
- Sample the initial proposal of slope and intercept using RANSAC
- Conduct RANSAC-based MH update for 5 iterations

```
function ransac_inference(xs, ys, observations)
    (slope, intercept) = ransac(xs, ys, RANSACParams(10, 3, 1.))
    slope_intercept_init = choicemap()
    slope_intercept_init[:slope] = slope
    slope_intercept_init[:intercept] = intercept
    (tr, _) = generate(regression_with_outliers, (xs,), merge(observations, slope_intercept_init))
    for iter=1:5
        tr = ransac_update(tr)
    end
    tr
end
```

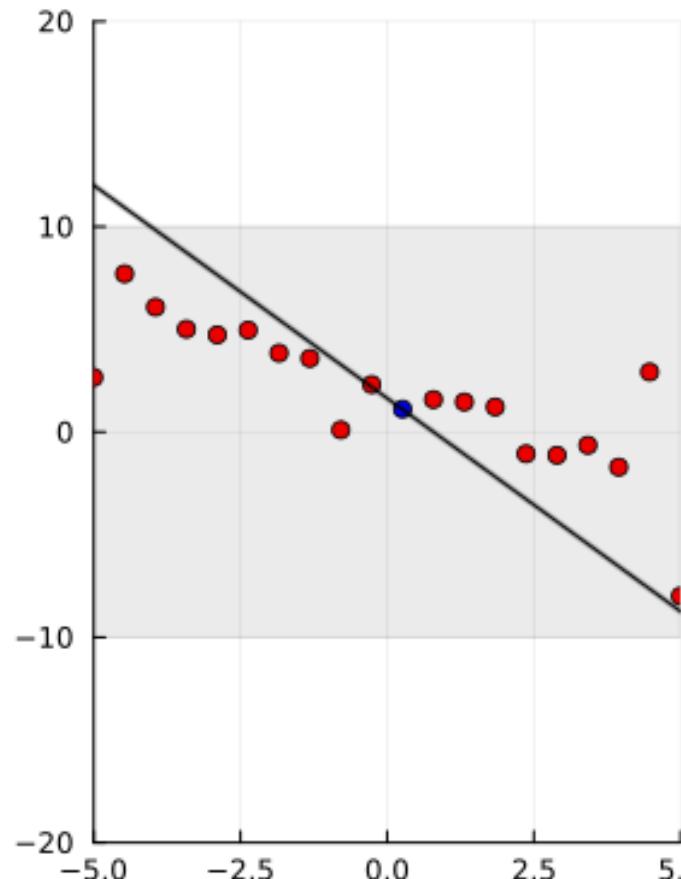
RANSAC-based MCMC inference

RANSAC Kernel

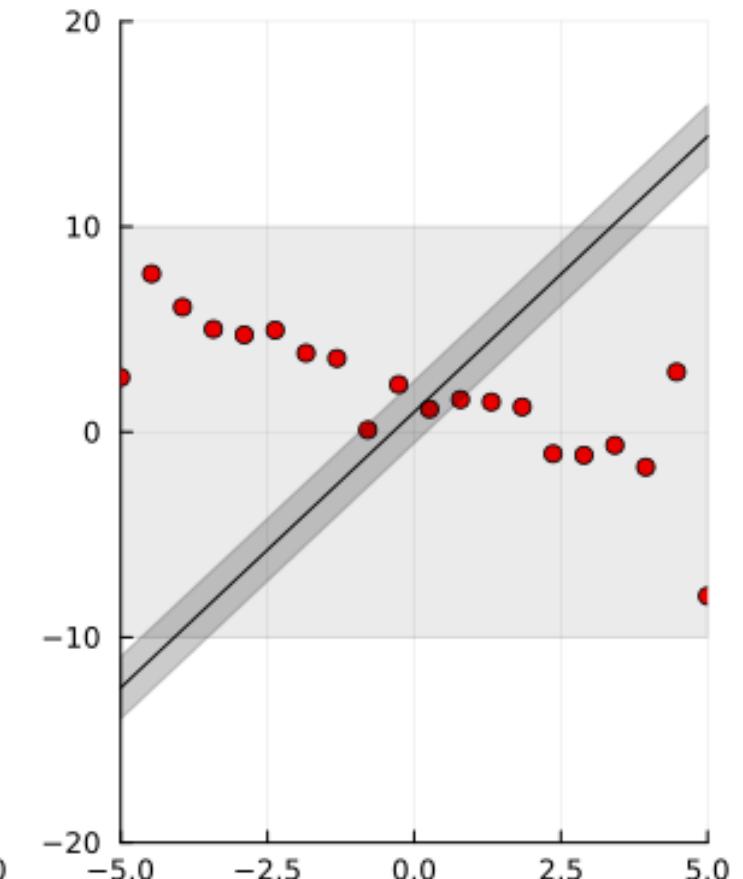
Iteration 1



Drift Kernel (Iter 1)



Resim Kernel (Iter 1)



Finds a reasonable proposal much earlier thanks to the heuristics (data-driven proposal)

Block resimulation MH

```
1.343702 seconds (11.47 M allocations: 676.832 MiB, 4.73% gc time, 29.51% compilation time)
0.802108 seconds (11.06 M allocations: 648.812 MiB, 5.46% gc time)
0.729552 seconds (11.06 M allocations: 648.812 MiB, 5.89% gc time)
0.787784 seconds (11.06 M allocations: 648.812 MiB, 5.40% gc time)
0.955552 seconds (11.06 M allocations: 648.812 MiB, 4.78% gc time)
0.891746 seconds (11.06 M allocations: 648.812 MiB, 5.89% gc time)
0.920517 seconds (11.06 M allocations: 648.812 MiB, 6.03% gc time)
0.909784 seconds (11.06 M allocations: 648.812 MiB, 5.95% gc time)
0.839699 seconds (11.06 M allocations: 648.812 MiB, 5.54% gc time)
0.713001 seconds (11.06 M allocations: 648.812 MiB, 5.84% gc time)
Log probability: -50.78536994535881
```

Gaussian drift MH

```
1.051416 seconds (12.01 M allocations: 699.933 MiB, 5.93% gc time, 1.19% compilation time)
1.161278 seconds (12.00 M allocations: 699.043 MiB, 5.36% gc time)
1.032212 seconds (12.00 M allocations: 699.043 MiB, 5.30% gc time)
0.732046 seconds (12.00 M allocations: 699.043 MiB, 6.20% gc time)
0.967897 seconds (12.00 M allocations: 699.043 MiB, 5.54% gc time)
0.947385 seconds (12.00 M allocations: 699.043 MiB, 5.59% gc time)
1.106322 seconds (12.00 M allocations: 699.043 MiB, 5.77% gc time)
1.220370 seconds (12.00 M allocations: 699.043 MiB, 4.90% gc time)
0.815143 seconds (12.00 M allocations: 699.043 MiB, 5.54% gc time)
0.778490 seconds (12.00 M allocations: 699.043 MiB, 6.39% gc time)
Log probability: -44.24115015595737
```

RANSAC proposal MH

```
0.930152 seconds (3.70 M allocations: 230.803 MiB, 3.76% gc time, 76.41% compilation time)
0.195143 seconds (2.41 M allocations: 144.268 MiB, 5.27% gc time)
0.179629 seconds (2.41 M allocations: 144.268 MiB, 9.60% gc time)
0.160974 seconds (2.41 M allocations: 144.268 MiB, 5.74% gc time)
0.169705 seconds (2.41 M allocations: 144.268 MiB, 7.40% gc time)
0.187931 seconds (2.41 M allocations: 144.268 MiB, 5.79% gc time)
0.193349 seconds (2.41 M allocations: 144.268 MiB, 8.71% gc time)
0.203889 seconds (2.41 M allocations: 144.268 MiB, 5.47% gc time)
0.194845 seconds (2.41 M allocations: 144.268 MiB, 6.39% gc time)
0.179412 seconds (2.41 M allocations: 144.268 MiB, 7.83% gc time)
Log probability: -42.80176245244315
```

- Better fitness
- Faster inference

No data-driven initial proposal?

```
function ransac_inference(xs, ys, observations)
    (slope, intercept) = ransac(xs, ys, RANSACParams(10, 3, 1.))
    slope_intercept_init = choicemap()
    slope_intercept_init[:slope] = slope
    slope_intercept_init[:intercept] = intercept
    (tr, _) = generate(regression_with_outliers, (xs,), merge(observations, slope_intercept_init))
    for iter=1:5
        tr = ransac_update(tr)
    end
    tr
end
```

Log probability: -42.80176245244315

```
function ransac_inference_no_init(xs, ys, observations)
    (tr, _) = generate(regression_with_outliers, (xs,), observations)
    for iter=1:5
        tr = ransac_update(tr)
    end
    tr
end
```

Log probability: -43.87146997454299

Summary: MCMC in Gen

- Define a generative model
- Generate a trace by running the generative model

```
(trace::U, weight) = generate(gen_fn::GenerativeFunction{T,U},  
args::Tuple, constraints::ChoiceMap)
```

- Update *selected* variables in the trace (a sweep: update all variables)

```
(new_trace, accepted) = mh(trace, selection::Selection; ...)
```

```
(new_trace, accepted) = mh(trace,  
proposal::GenerativeFunction, proposal_args::Tuple;  
selection::Selection; ...)
```

Try it out yourself

- Go over the jupyter notebook “MCMC in Gen.ipynb”

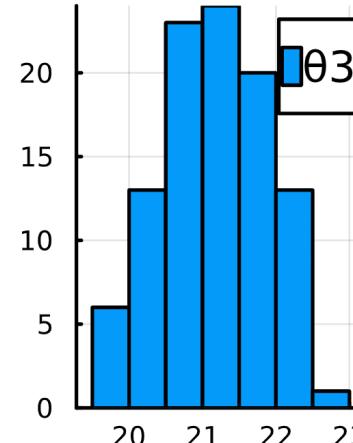
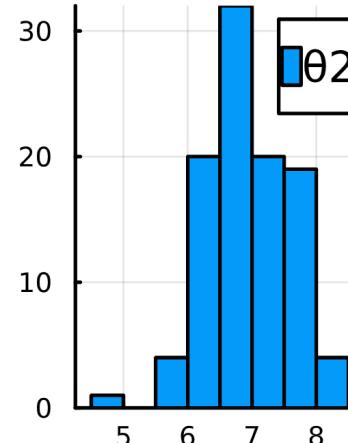
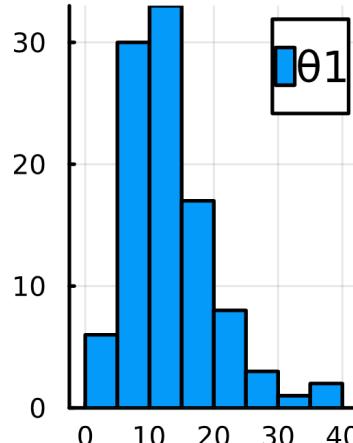
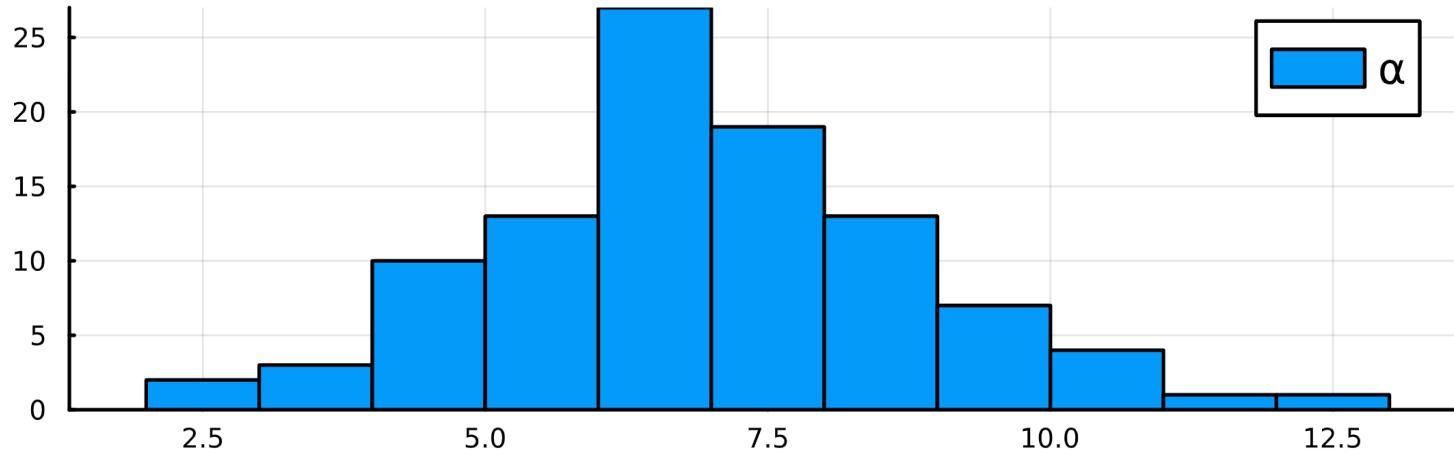
Problem set 2 release

- Due by March 26, end of day
- Q1: importance sampling vs MCMC in Gen
 - Implement both
 - MCMC requires a random walk MH update
 - Hint: uniformly sample from a range centered around the previous value
 - Evaluation: average log probability after several runs

A successful inference result based on importance sampling

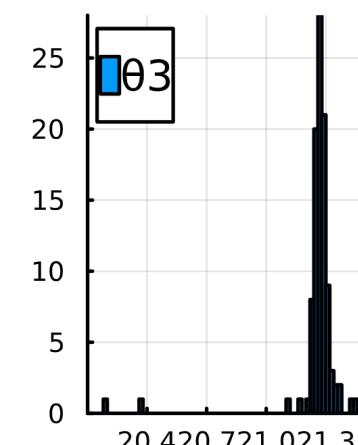
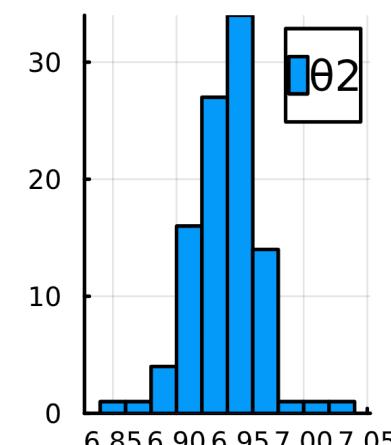
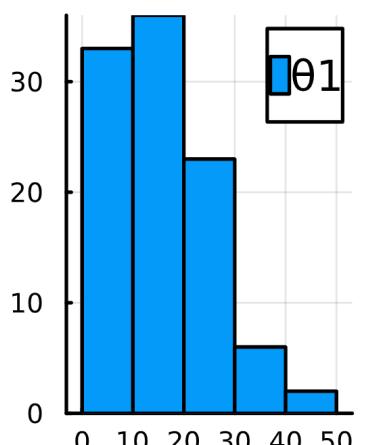
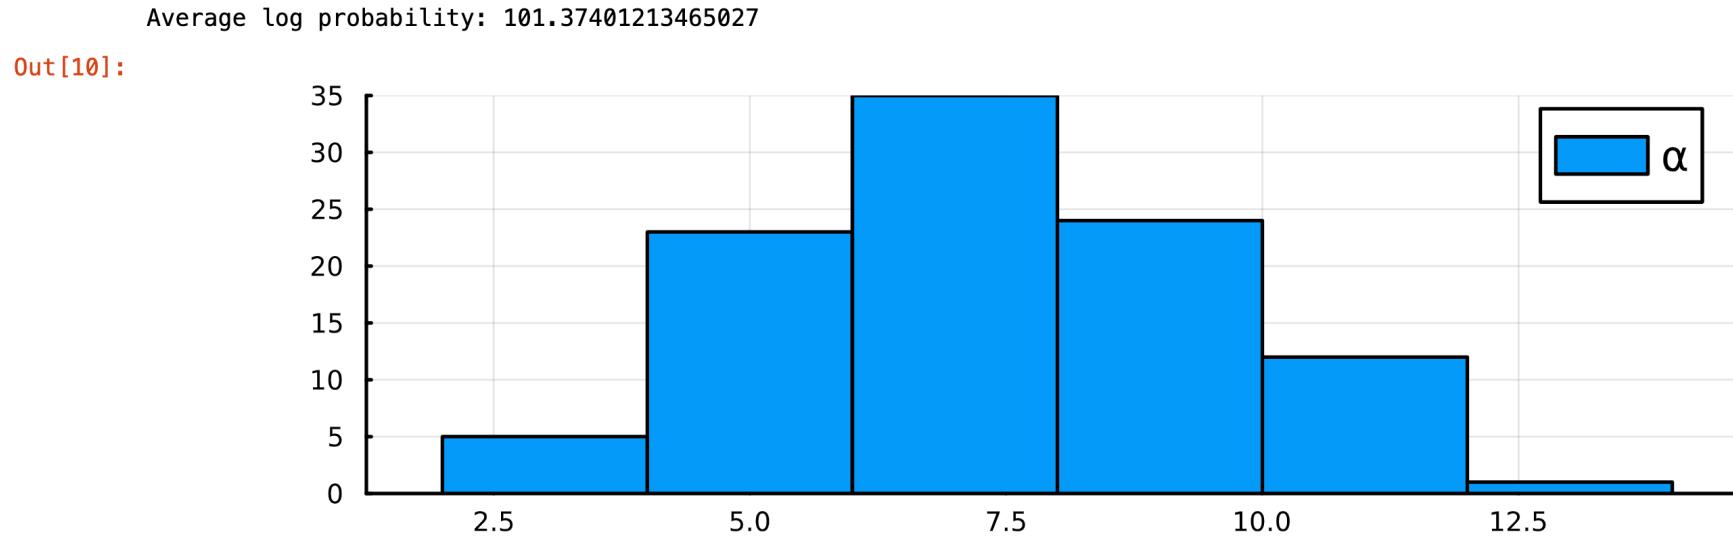
Average log probability: 88.14669239667002

Out[7]:



A successful inference result based on MCMC

You won't get the exact same result, but the average log prob for MCMC should be at least around 100



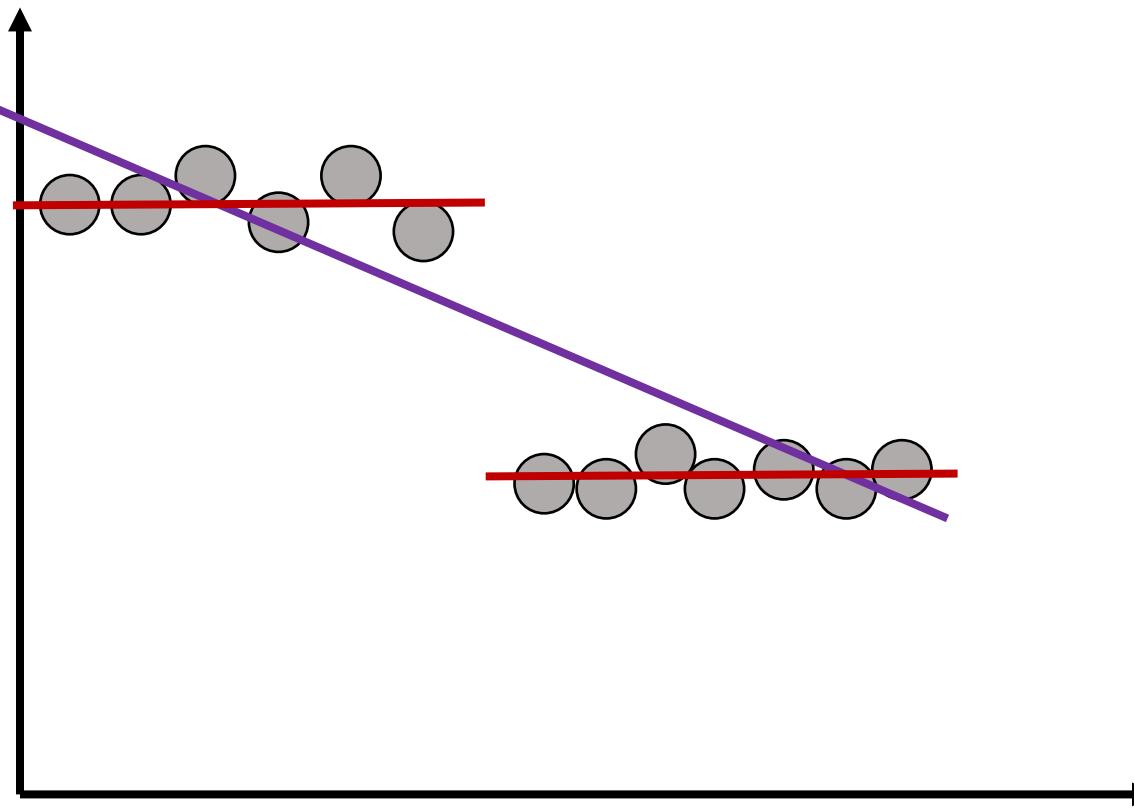
Summary

- MCMC provides an iterative update of hypotheses
- Start from an initial hypothesis (from a prior or heuristics-based/data-driven proposal)
- Update the hypothesis via Gibbs sampling or Metropolis-Hastings
- Smart proposals will help you converge to the true posterior distribution faster (mix well) – customized proposals in Gen

- Currently, we assume a fixed set parameters for each hypothesis
- How can we generalize this to hypotheses with different parameters / dimensions?

An example of hypotheses with different dimensions

- An unknown number of lines



An example of hypotheses with different dimensions

- Consider a generative model of two pieces of observed data, y_1 and y_2

```
@gen function model()
    if (:{z} ~ bernoulli(0.5))
        m1 = (:{m1} ~ gamma(1, 1))
        m2 = (:{m2} ~ gamma(1, 1))
    else
        m = (:{m} ~ gamma(1, 1))
        (m1, m2) = (m, m)
    end
    {:{y1}} ~ normal(m1, 0.1)
    {:{y2}} ~ normal(m2, 0.1)
end
```

Two models:

- 1) Same normal dist. with a mean of m
- 2) Two normal dist., with two means, m_1 , m_2

Selection MH

- Selection kernel (select one of the model types):

```
select_mh_structure_kernel(trace) = mh(trace, select(:z))[1]
```

- Mean proposal kernel (proposal mean(s) given model type):

```
@gen function fixed_structure_proposal(trace)
    if trace[:z]
        {:m1} ~ normal(trace[:m1], 0.1)
        {:m2} ~ normal(trace[:m2], 0.1)
    else
        {:m} ~ normal(trace[:m], 0.1)
    end
end

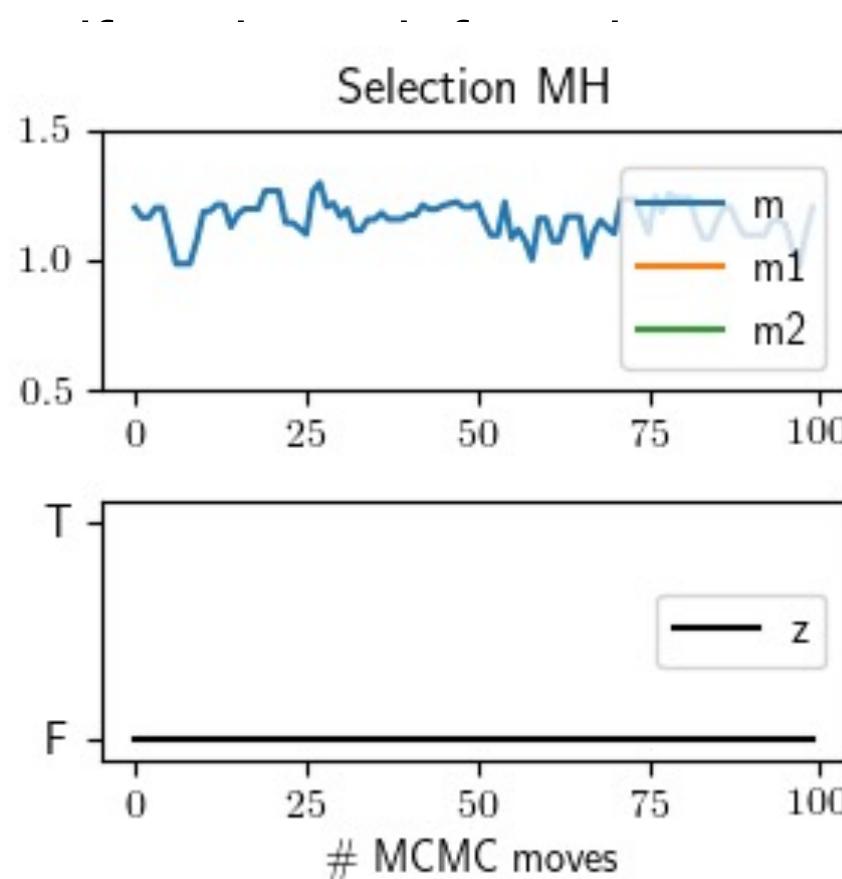
fixed_structure_kernel(trace) = mh(trace, fixed_structure_proposal,())[1]
```

- MH updates using the two kernels:

```
(y1, y2) = (1.0, 1.3)
trace, = generate(model, (), choicemap(:y1, y1), (:y2, y2), (:z, false), (:m, 1.2))
for iter=1:100
    trace = select_mh_structure_kernel(trace)
    trace = fixed_structure_kernel(trace)
end
```

Problem: inefficient proposal sampling

- When switching from the model with a single mean to the model with two means, the values of the new addresses `:m1` and `:m2` will be proposed from the prior distribution (again).
- This is wasteful, since we expect the values for `:m1` and `:m2` to be true when proposing them.
- This means it will take longer to find the posterior probability.



ate value for `:m`, we expect the same value. The same is true when proposing in the opposite direction. This means it will take longer to find an accurate estimate of the posterior probability.

Reversible Jump MCMC (RJMCMC)

- Standard MH: a fixed number of variables
- Different hypotheses have different dimensions
- Reversible jump MCMC is just MH, extended to moves between different hypotheses with varying dimensions, which is common for Bayesian model selection
- Key: to ensure that MH converges to the correct stationary distribution, the proposal distributions (jumps between hypotheses) must satisfy detailed balance:

$$p(h|D)q(h'|h)a(h'|h) = p(h'|D)q(h|h')a(h|h')$$

Reversible proposals

Geometric mean & ratio

$(m_1, m_2) \rightarrow (m, u)$

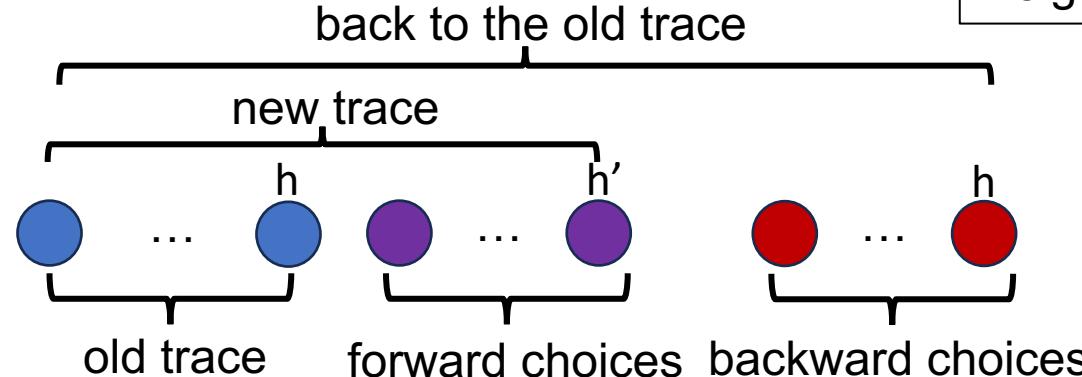
```
function merge_means(m1, m2)
    m = sqrt(m1 * m2)
    u = m1 / (m1 + m2)
    (m, u)
end
```

$(m, u) \rightarrow (m_1, m_2)$

```
function split_mean(m, u)
    m1 = m * sqrt(u / (1 - u))
    m2 = m * sqrt((1 - u) / u)
    (m1, m2)
end
```

RJMCMC in Gen

- Involution functions: allow reversed jumps $h \leftarrow h'$



Inputs:

old_trace, forward_choices

Outputs:

new_trace, backward_choices,
weight = LL(new_trace) – LL(old_trace)

- One way to write involution function: write a Julia function $f()$

$h \rightarrow h'$ (choices that lead to h' from h) trace of h

```
forward_choices = get_choices(simulate(proposal_randomness_func, (old_trace,)))
```

```
new_trace, backward_choices, = f(old_trace, forward_choices, ...)
```

$h \rightarrow h'$ (choices that lead to h' from h)

trace of h'

```
back_to_the_first_trace_maybe, backward_backward_choices, = f(new_trace, backward_choices, ...)
```

trace of h

$h \rightarrow h'$ (choices that lead to h' from h)

$h' \rightarrow h$ (choices that lead to h from h')

`backward_backward_choices == forward_choices`

`back_to_the_first_trace_maybe == old_trace`

RJMCMC in Gen

- A simpler way to write involution function:
- Write a customized forward proposal function that only sample auxiliary variables that are not part of the generative model (e.g., u)
- Then write a transform function that implement a reversible kernel based on the model variables and the auxiliary variables. This kernel can change model variables in forward and backward directions based on the auxiliary variables

Split/merge kernel

Split a variable or merge two variables

```
@gen function split_merge_proposal(trace)
    if trace[:z]
        # currently two segments, switch to one
    else
        # currently one segment, switch to two
        {:u} ~ uniform_continuous(0, 1)
    end
end
```

Only propose auxiliary variables (variable not part of the generative model)

Split/merge involution function

A trace transform function:

Transforms read the value of random choices from input trace(s) and write values of random choices to output trace(s)

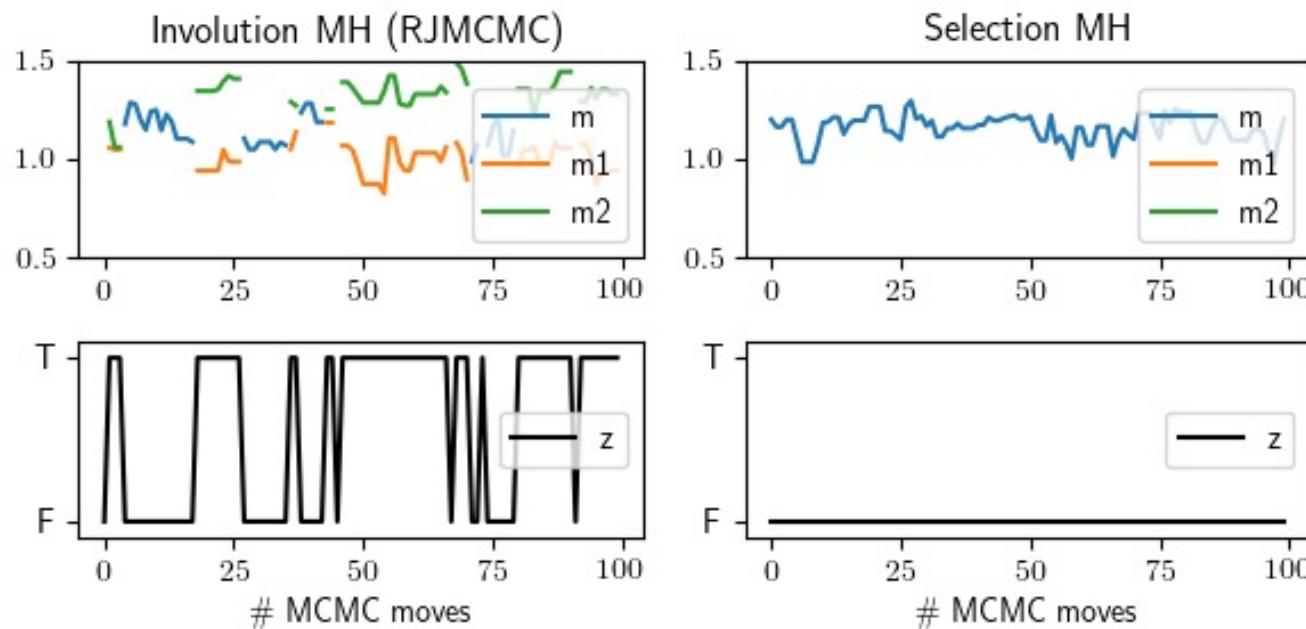
```
@transform split_merge_involution (model_in, aux_in) to (model_out, aux_out) begin
    if @read(model_in[:z], :discrete)

        # currently two means, switch to one
        @write(model_out[:z], false, :discrete)
        m1 = @read(model_in[:m1], :continuous)
        m2 = @read(model_in[:m2], :continuous)
        (m, u) = merge_means(m1, m2)
        @write(model_out[:m], m, :continuous)
        @write(aux_out[:u], u, :continuous)
    else

        # currently one mean, switch to two
        @write(model_out[:z], true, :discrete)
        m = @read(model_in[:m], :continuous)
        u = @read(aux_in[:u], :continuous)
        (m1, m2) = split_mean(m, u)
        @write(model_out[:m1], m1, :continuous)
        @write(model_out[:m2], m2, :continuous)
    end
end
```

RJMCMC based on split/merge kernel

```
split_merge_kernel(trace) = mh(trace, split_merge_proposal, (), split_merge_involution)[1]
(y1, y2) = (1.0, 1.3)
trace, = generate(model, (), choicemap(:y1, y1), (:y2, y2), (:z, false), (:m, 1.))
for iter=1:100
    trace = split_merge_kernel(trace)
    trace = fixed_structure_kernel(trace)
    println(trace[:z])
end
```



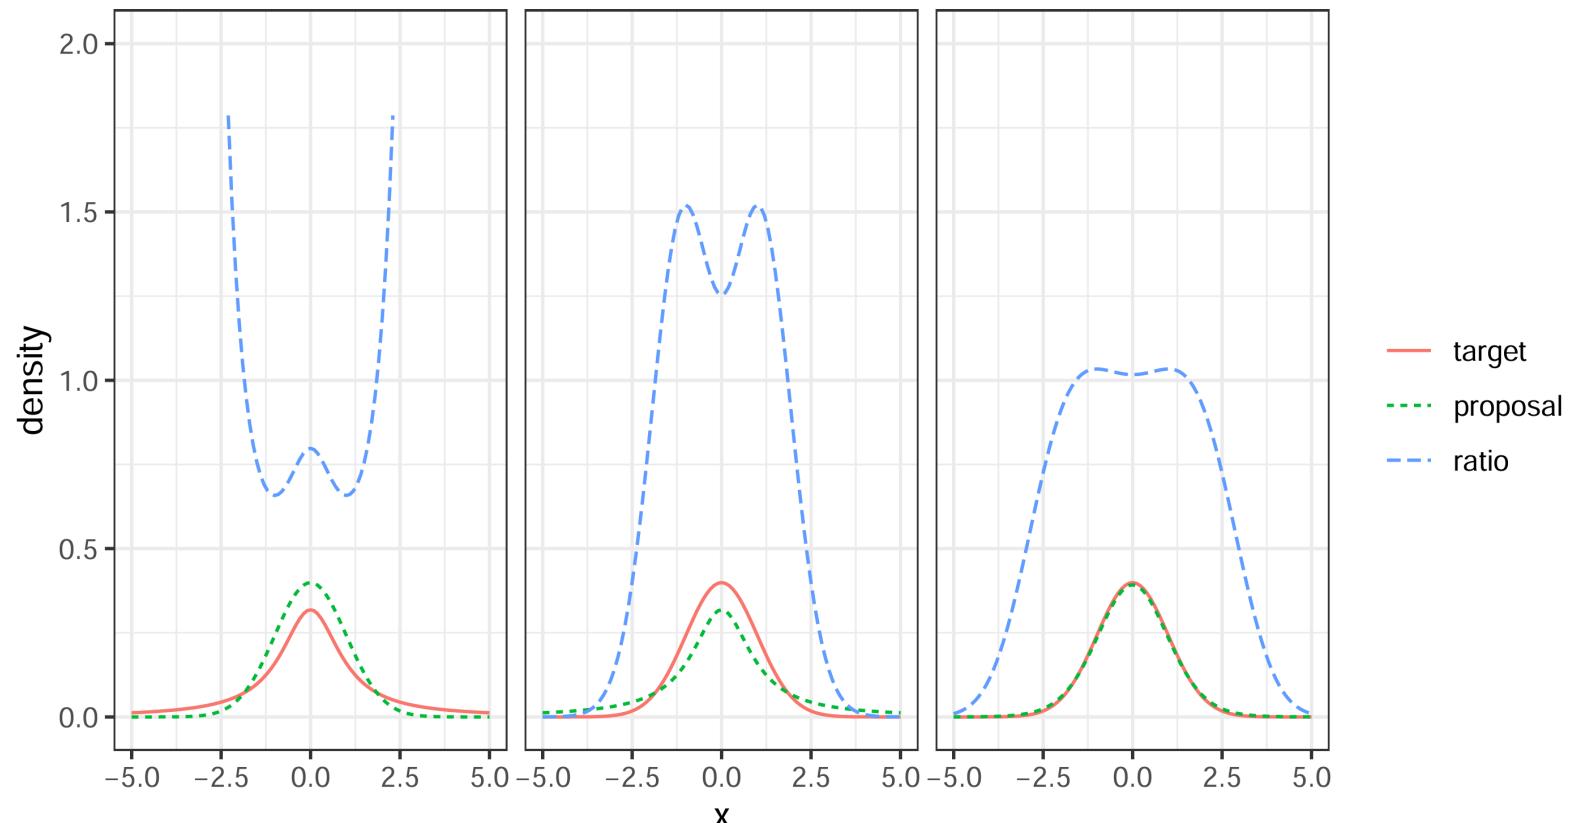
Sampling-based inference algorithms

- Monte Carlo methos:
 - Rejection sampling
 - Importance sampling
 - Markov chain Monte Carlo (MCMC)
 - Sequential Monte Carlo (SMC)

How to use Gen to implement sampling-based inference algorithms?

Recap: Importance sampling

- **Goal:** approximate a *target* distribution $P(x)$, which is hard to sample from
- **Main idea:** Sample a set of particles from a simple *proposal* distribution $q(x)$ (e.g., a uniform distribution) and *weight* them to approximate the distribution
- For each particle i
 - Sample: $x_i \sim q(x)$
 - Weight: $w_i = \frac{p(x_i)}{q(x_i)}$
 - (x_i, w_i)



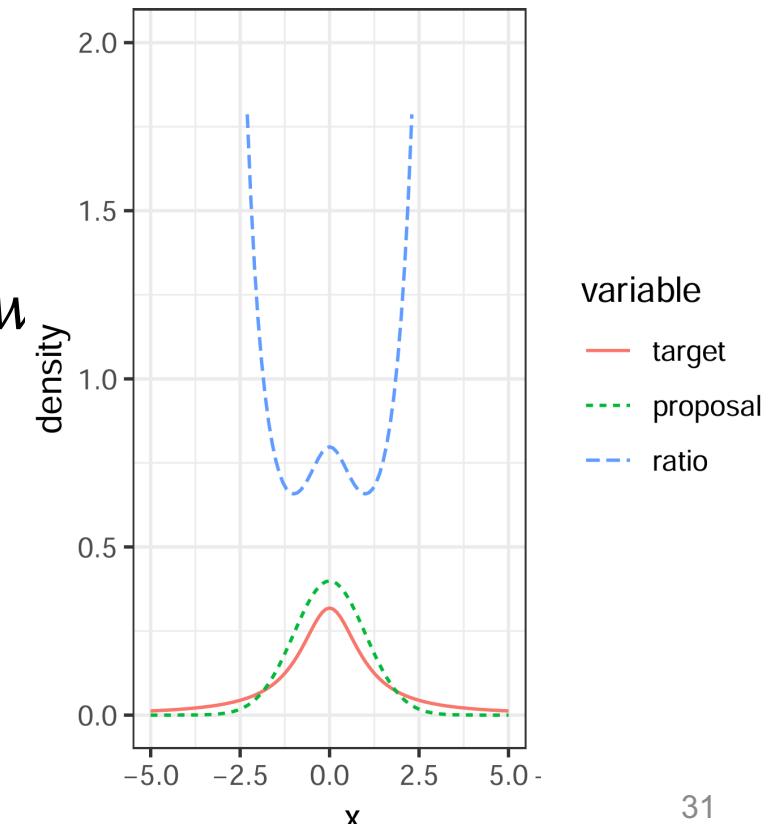
Recap: Importance resampling

- **Goal:** sample from the approximated distribution based on the particles and their weights
- Normalize the weights of particles (probability of each particle assuming the whole population is the sampled particles):

$$w_i = \frac{w_i}{\sum_j w_j}$$

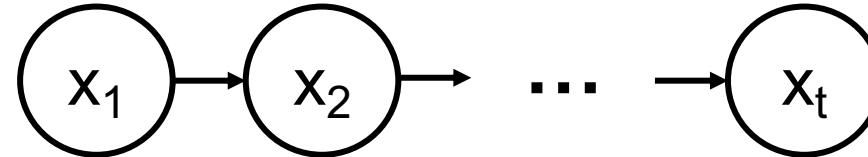
- For each new sample x ,
 - Sample $j \in \{1, 2, \dots, K\}$, from with probability w_j
 - The new sample $x = x_j$
 - Set weight $w_i = 1/K$

Limit: proposal distribution needs to be reasonably close to the target distribution



Sequential data

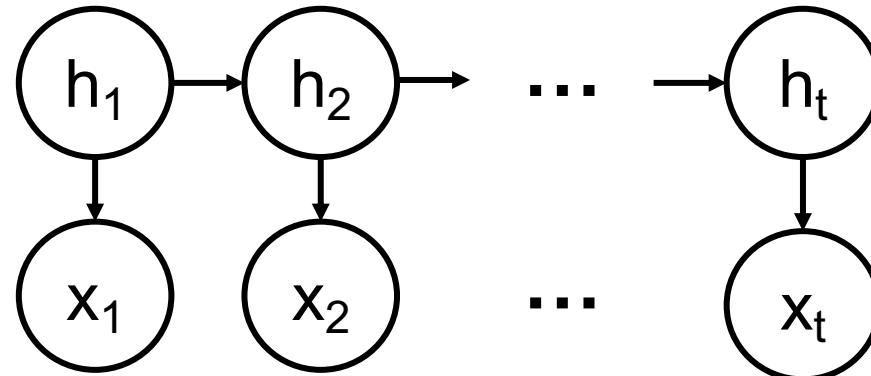
- Sample sequential data: $\{x_1, x_2, \dots, x_t\} \sim p(x_1, x_2, \dots, x_t)$
- Markov chain: $p(x_1, x_2, \dots, x_t) = p(x_1) \prod_{\tau=2}^t p(x_\tau | x_{\tau-1})$



- Sample the initial variable: $x_1 \sim p(x_1)$
- Conditional sampling for other time steps: $x_t \sim p(x_t | x_{t-1}, \dots, x_1)$
- **Basic idea for Sequential Monte Carlo (SMC):** For every step, we can approximate the true target distributions (the initial distribution and the conditional distribution at every step) using importance resampling

SMC for Bayesian inference

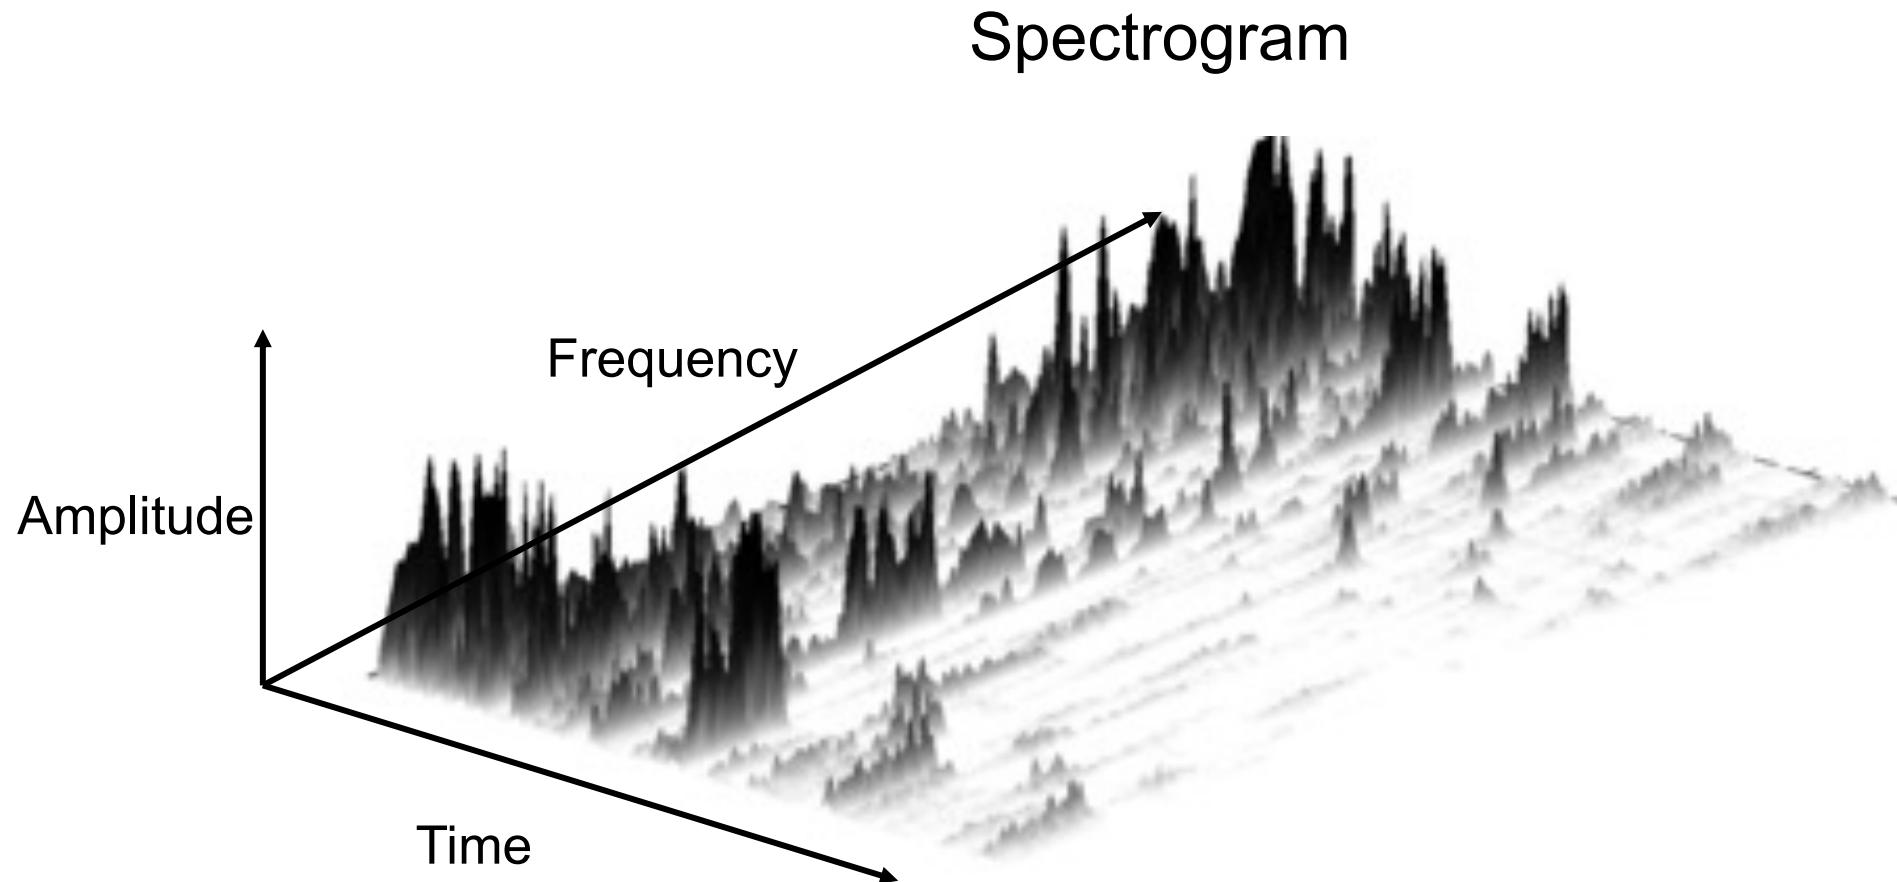
- Sample hypotheses for sequential data:
- $\{h_1, h_2, \dots, h_t\} \sim p(h_1, h_2, \dots, h_t | x_1, x_2, \dots, x_t)$
- A Bayesian network:



- Sample the initial variable: $h_1 \sim p(h_1 | x_1)$
- Conditional sampling for other time steps: $h_t \sim p(h_t | h_{t-1}, x_1, \dots, x_t)$

SMC in Gen

- Example 1: Music notes



SMC in Gen

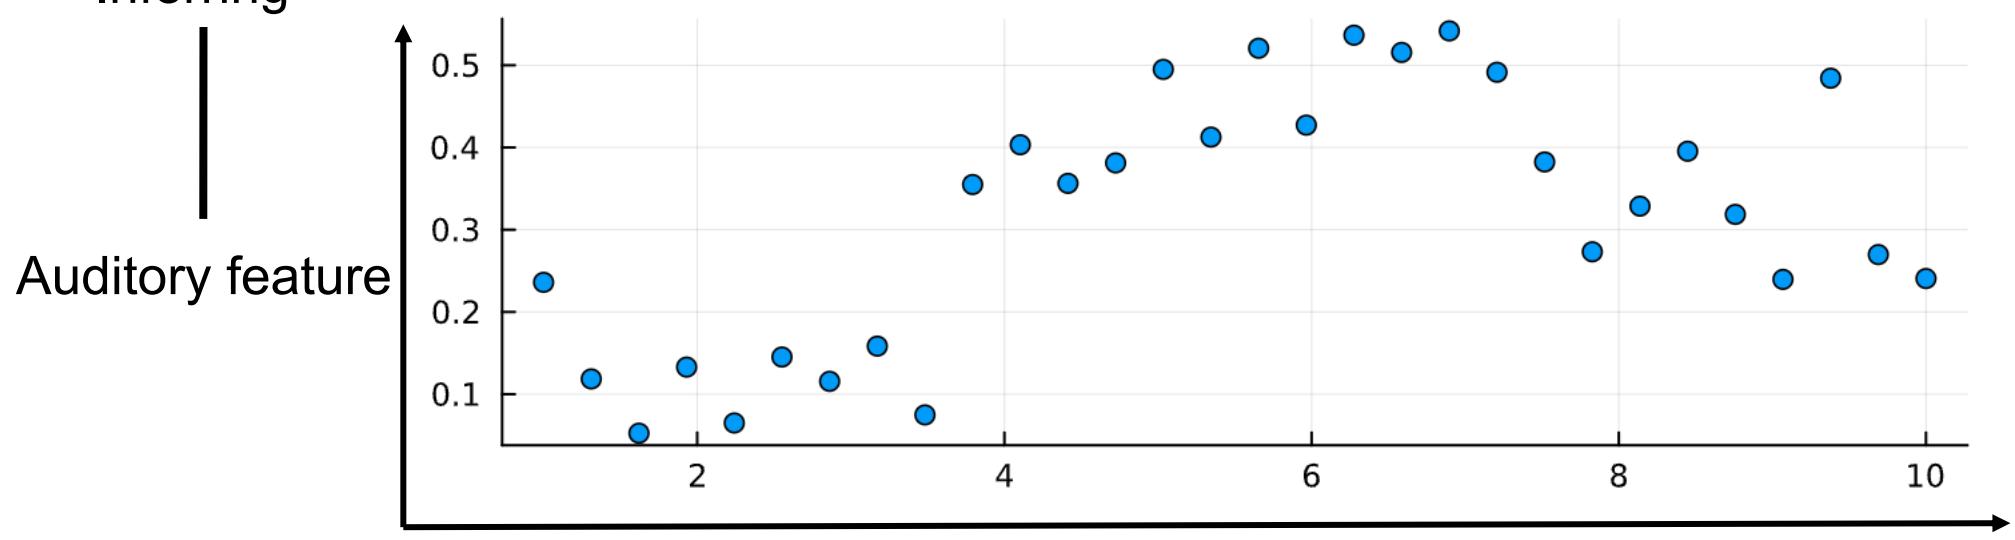
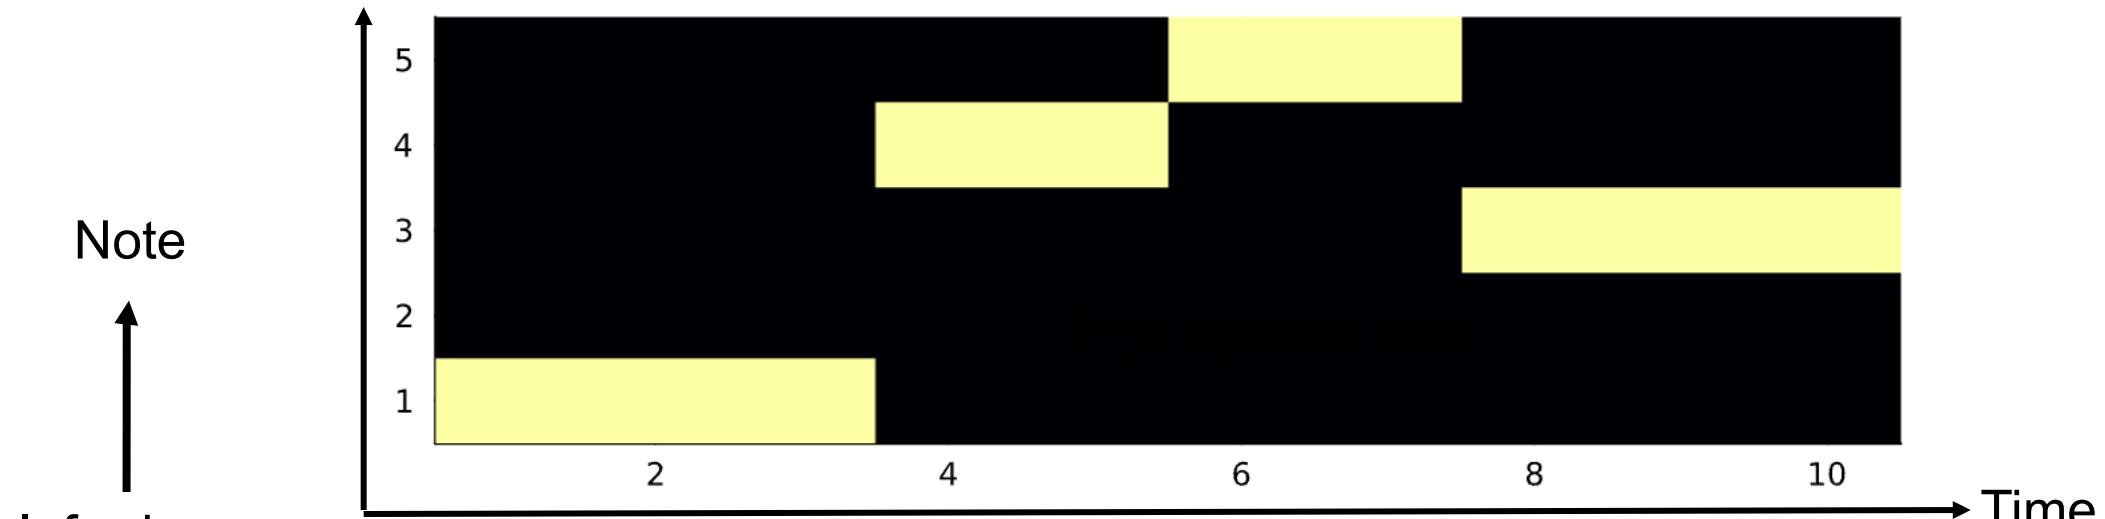
- Example 1: Music notes

$$p(h_1, h_2, \dots, h_t | x_1, x_2, \dots, x_t)$$

A simplified model

Why SMC instead of MCMC?

Temporal dependency



SMC

- Step 1: Initialization
 - Set $t = 1$
 - For $i \in \{1, \dots, K\}$, sample $h_1^{(i)} \sim q(h_1)$ initial proposals
 - Weight $w_1^{(i)} = \frac{p(x_1 | h_1^{(i)}) p(h_1^{(i)})}{q(h_1^{(i)})}$
 - Normalize all weights \rightarrow particles $\{(h_1^{(i)}, w_1^{(i)})\}_{i=1}^K$
 - Step 2: Resampling at step t
 - If the particles are not diverse enough, resample the particles, and set $w_t^{(i)} = \frac{1}{K}$ (importance resampling)
 - Step 3: Sampling at step $t + 1$
 - Set $t = t + 1$; particles at previous step $\{(h_{1:t-1}^{(i)}, w_{t-1}^{(i)})\}_{i=1}^K$
 - For $i \in \{1, \dots, K\}$, sample $h_t^{(i)} \sim p(h_t | h_{t-1}^{(i)})$ conditional distribution
 - Weight $w_t^{(i)} = w_{t-1}^{(i)} p(x_t | h_t^{(i)})$ reweighted by the current likelihood
 - Normalize all weights \rightarrow particles $\{(h_{1:t}^{(i)}, w_t^{(i)})\}_{i=1}^K$
- hypotheses at all previous steps
-

When to resample?

- Effective sample size (ESS)

$$ESS = \left(\sum_{i=1}^K (w^{(j)})^2 \right)^{-1}$$

- Large sample size: particles share similar weights
- Small sample size: only a small number of particles have large weights

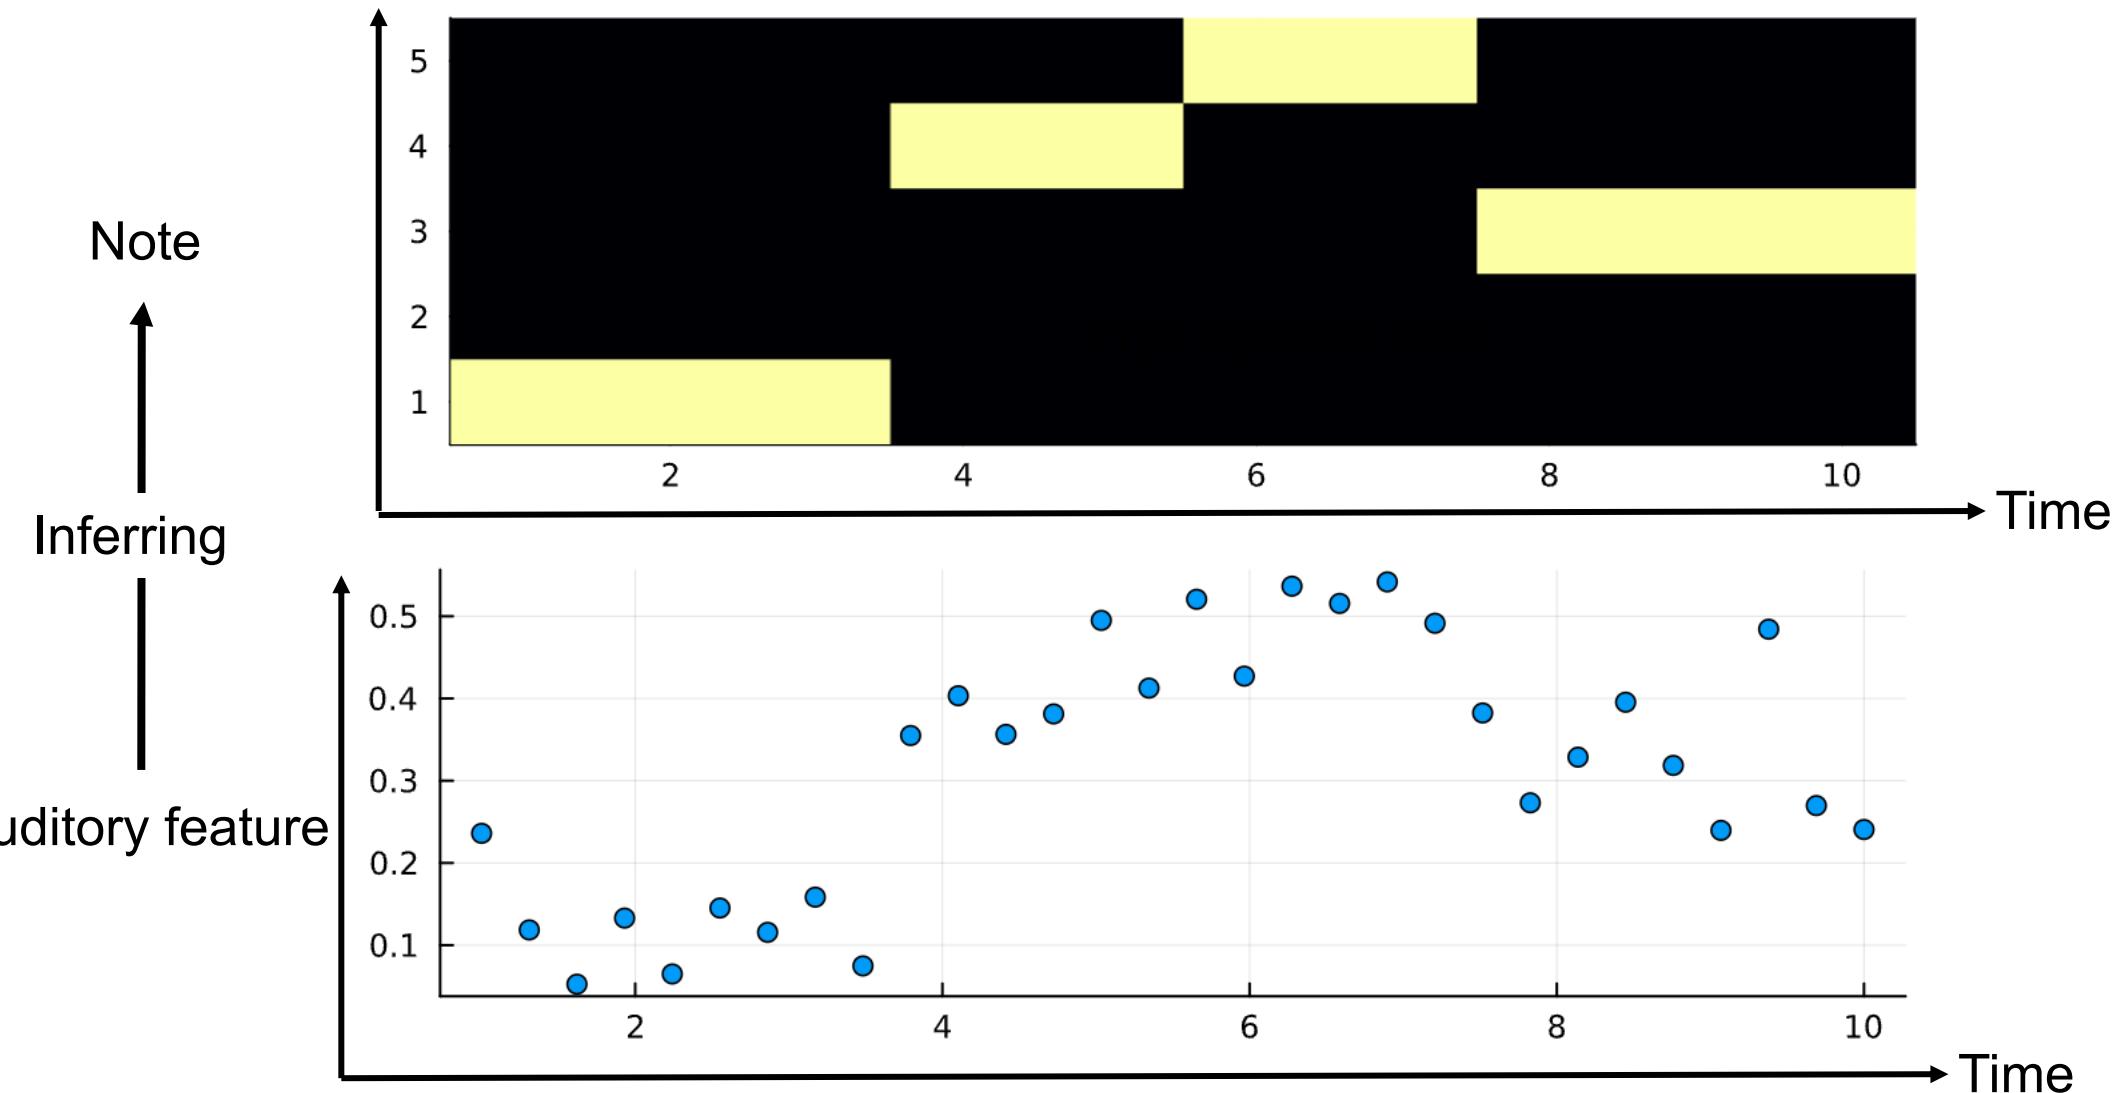
SMC

- Step 1: Initialization
 - Set $t = 1$
 - For $i \in \{1, \dots, K\}$, sample $h_1^{(i)} \sim q(h_1)$ initial proposals
 - Weight $w_1^{(i)} = \frac{p(x_1 | h_1^{(i)}) p(h_1^{(i)})}{q(h_1^{(i)})}$
 - Normalize all weights \rightarrow particles $\{(h_1^{(i)}, w_1^{(i)})\}_{i=1}^K$
- Step 2: Resampling at step t
 - If **ESS < threshold**, resample the particles, and set $w_t^{(i)} = \frac{1}{K}$ (importance resampling)
- Step 3: Sampling at step $t + 1$
 - Set $t = t + 1$; particles at previous step $\{(h_{1:t-1}^{(i)}, w_{t-1}^{(i)})\}_{i=1}^K$
 - For $i \in \{1, \dots, K\}$, sample $h_t^{(i)} \sim p(h_t | h_{t-1}^{(i)})$ conditional distribution
 - Weight $w_t^{(i)} = w_{t-1}^{(i)} p(x_t | h_t^{(i)})$ reweighted by the current likelihood
 - Normalize all weights \rightarrow particles $\{(h_{1:t}^{(i)}, w_t^{(i)})\}_{i=1}^K$

SMC in Gen (particle filtering)

- Example 1: Music notes

A simplified model



Generative model – transitions between notes $p(h_t|h_{t-1})$

- 5 musical notes, a composition is defined by a set of stochastic transitions between notes
- With a probability of 0.6, we stay on the current note; with a probability of 0.1, we transition to one of the new notes

```
function transition_prob(a, b)
    a == b ? 0.6 : 0.1
end

# setup transition probabilities
function transition_prob_matrix()
    note_transitions = Array{Float64}(undef, 5, 5)
    for i=1:5
        [note_transitions[i, j] = transition_prob(i,j) for j=1:5]
    end
    return note_transitions
end

# make this variable global (makes some of the coding more straightforward)
global note_transitions = transition_prob_matrix()
```

Generative model – notes & auditory features $p(x_{1:t}, h_{1:t})$

```
@gen function musical_notes(K::Int)
    # sample an initial "note"
    current_note = {::notes => 1 => :θ} ~ uniform_discrete(1, 5)

    # "play" it: projecting it to "auditory features" of length 3
    note_play_length = 3
    mu = current_note * 0.1
    {::data => 1 => :y} ~ broadcasted_normal(repeat([mu], note_play_length), repeat([0.05], note_play_length))

    # keep going for K notes
    for k=1:K
        current_note = {::notes => k+1 => :θ} ~ categorical(note_transitions[current_note, :])
        mu = current_note * 0.1
        {::data => k+1 => :y} ~ broadcasted_normal(repeat([mu], note_play_length), repeat([0.05], note_play_length))
    end
end
```

See jupyter notebook